

# Power simulation for linear mixed-effects models

Nora Wickelmaier

February 6, 2023

# Power analysis by simulation

Why simulation?

- Simulation is at the heart of statistical inference
- Inference: Compare the data with the output of a statistical model
- If data look different from model output, reject model (or its assumptions)
- Simulation forces us to **specify a data model** and to attach meaning to its components
- Model should not be totally unrealistic for those aspects of the world we want to learn about

## Specify the model including the effect of interest

(1) Choose statistical model according to its assumptions

- Binomial test  $\rightarrow$  binomial distribution  $\rightarrow$  `rbinom()`
- t test  $\rightarrow$  normal distribution  $\rightarrow$  `rnorm()`
- ...

(2) Fix unknown quantities

- Standard deviations, correlations, ...

(3) Specify the effect of interest

- Smallest effect size of interest (SESOI)

# Estimating power with simulation

## Pseudo code

```
Set sample size
replicate
{
  Draw sample from model with minimal relevant effect
  Test null hypothesis
}
Determine proportion of significant results
```

## Sample size calculation

- Adjust  $n$  until desired power (0.8 or 0.95) is reached
- To be on the safe side, assume higher variation, less (or more) correlation, and smaller interesting effects (what results can we expect, if ...)

# Generalized linear models

- A generalized linear model is defined by

$$g(E(y)) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k,$$


where  $g()$  is the link function that links the mean to the linear predictor. The response  $y$  is assumed to be independent and to follow a distribution from the exponential family

- In R, a GLM is fitted by

```
glm(y ~ x1 + x2 + ... + xk, family(link), data)
```

## Binomial regression

- Logit or probit models are special cases of GLMs for binomial response variables
- Artificial example: congenital eye disease



Logit model

$$\log \frac{p}{1-p} = \beta_0 + \beta_1 AGE$$

Probit model

$$\Phi^{-1}(p) = \beta_0 + \beta_1 AGE$$

## Logit model / logistic regression

- We want to model the probability of  $y$  with the logistic function

$$p(y) = \frac{1}{1 + e^{-y}} \quad \text{with } y \sim \text{Binom}(n, p)$$

- How do we get the logit model  $\log \frac{p}{1-p} = \beta_0 + \beta_1 x$  from that?

$$\begin{aligned} \log \left( \frac{p}{1-p} \right) &= \log(p) - \log(1-p) = \log \left( \frac{1}{1 + e^{-y}} \right) - \log \left( 1 - \frac{1}{1 + e^{-y}} \right) \\ &= \log(1) - \log(1 + e^{-y}) - \log(e^{-y}) + \log(1 + e^{-y}) \\ &= -\log(e^{-y}) \\ &= y := \beta_0 + \beta_1 x \end{aligned}$$

with  $1 - \frac{1}{1+e^{-y}} = \frac{e^{-y}}{1+e^{-y}}$

## Refresher: Logarithm rules

Product:  $\log(xy) = \log x + \log y$

Quotient:  $\log \frac{x}{y} = \log x - \log y$

Power:  $\log(x^p) = p \log x$

Root:  $\log \sqrt[p]{x} = \frac{\log x}{p}$



# Meta-information and Memory: Nicole's experiment

## Task:

Participants read 40 sentences that can be either true (30) or false (10) which are presented in different font colors

Base

The friends sit around a coffee table.

Low

The friends sit around a coffee table.

High

The friends sit around a coffee table.

## Study 1:

Recognition Memory: "Did you read this statement?"

DV – correct answer: correct recall

IVs – condition (control vs. high vs. low discriminability)  
– truthvalue (true vs. false)

## Study 2:

Source Memory: "Was this statement true, false, or new?"

DV – correct answer: true sentence classified as true

IVs – condition (high vs. low discriminability)  
– truthvalue (true vs. false)

# Meta-information and Memory: Nicole's experiment

## Hypotheses

### Study 1 Recognition Memory:

Significant interaction effect between truthvalue and discriminability: Recognition memory will be higher for true than false statements when a discriminability task is present

### Study 2 Source Memory:

Significant interaction effect between truthvalue and discriminability: Participants in high discriminability condition will more often classify statements correctly as true or false than in the low discriminability condition

# Meta-information and Memory: Nicole's experiment

## Hypotheses

### Study 1 Recognition Memory:

Significant interaction effect between truthvalue and discriminability: Recognition memory will be higher for true than false statements when a discriminability task is present

### Study 2 Source Memory:

Significant interaction effect between truthvalue and discriminability: Participants in high discriminability condition will more often classify statements correctly as true or false than in the low discriminability condition

→ We have the data of the first study and want to calculate the **number of participants** we need for the second study

# Meta-information and Memory: Nicole's experiment

## Model

- In order to test the interaction, we fit the following model:

$$\log \left( \frac{P(\text{correct})}{P(\text{false})} \right) = \beta_0 + \beta_1 \text{truthval} + \beta_2 \text{cond} + \beta_3 (\text{truthval} \times \text{cond}) + v_0 + \eta_0$$

with  $v_0 \sim N(0, \sigma_v)$  and  $\eta_0 \sim N(0, \sigma_\eta)$ , all random effects i.i.d.

- This model can be fitted in R with

```
lme4::glmer(correctAnswer ~ truthvalue * condition +  
             (1|participantID) + (1|itemID),  
             data, family = binomial)
```

# Meta-information and Memory: Nicole's experiment

## Model

- In order to test the interaction, we fit the following model:

$$\log \left( \frac{P(\text{correct})}{P(\text{false})} \right) = \beta_0 + \beta_1 \text{truthval} + \beta_2 \text{cond} + \beta_3 (\text{truthval} \times \text{cond}) + v_0 + \eta_0$$

with  $v_0 \sim N(0, \sigma_v)$  and  $\eta_0 \sim N(0, \sigma_\eta)$ , all random effects i.i.d.

- This model can be fitted in R with

```
lme4::glmer(correctAnswer ~ truthvalue * condition +  
              (1|participantID) + (1|itemID),  
              data, family = binomial)
```

→ How many parameters does this model have?

## Estimates from Study 1

Random effects	Variance	Std.Dev.
itemID	0.29	0.54
participantID	0.58	0.76

Fixed effects	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	1.62	0.15	11.15	0.00
truthvaluetrue	-0.02	0.12	-0.16	0.87
conditionhigh	-0.50	0.19	-2.59	0.01
conditionlow	-0.64	0.19	-3.35	0.00
truthvaluetrue:conditionhigh	0.39	0.16	2.44	0.01
truthvaluetrue:conditionlow	0.40	0.16	2.48	0.01

## Estimates from Study 1

Random effects	Variance	Std.Dev.
itemID	0.29	0.54
participantID	0.58	0.76

Fixed effects	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	1.62	0.15	11.15	0.00
truthvaluetrue	-0.02	0.12	-0.16	0.87
conditionhigh	-0.50	0.19	-2.59	0.01
conditionlow	-0.64	0.19	-3.35	0.00
truthvaluetrue:conditionhigh	0.39	0.16	2.44	0.01
truthvaluetrue:conditionlow	0.40	0.16	2.48	0.01

→ Parameters are on the logit scale

## Model predictions

truthvalue	condition	prob	logit
false	control	0.84	1.62
true	control	0.83	1.60
false	high	0.76	1.13
true	high	0.82	1.50
false	low	0.73	0.98
true	low	0.80	1.36

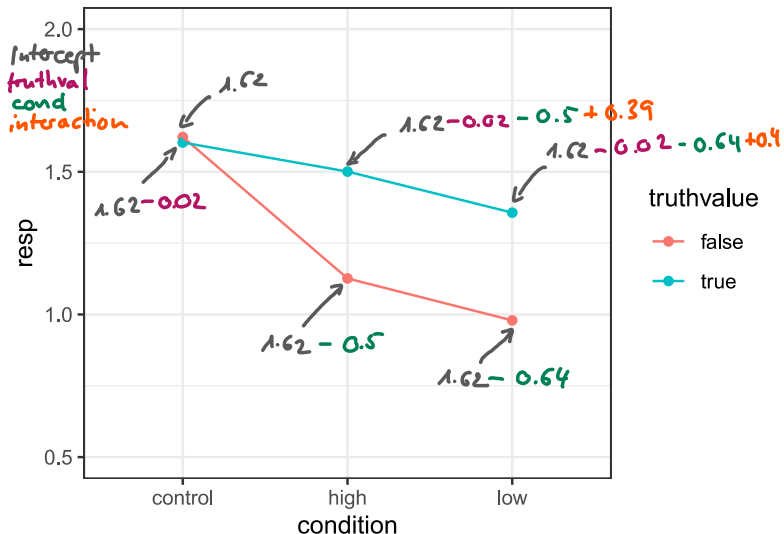


## Predictions on the logit scale

	est
(Intercept)	1.62
truthvaltrue	-0.02
condhigh	-0.50
condlow	-0.64
truthvaltrue:condhigh	0.39
truthvaltrue:condlow	0.40

Get predictions:

```
predict(glm1,  
        re.form = NA)
```

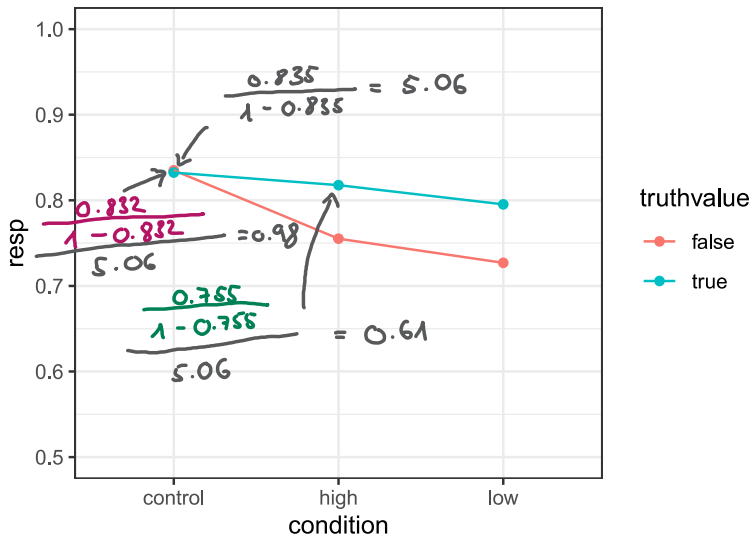


## Predictions on the probability scale

	OR
(Intercept)	5.06
truthvaltrue	0.98
condhigh	0.61
condlow	0.53
truthvaltrue:condhigh	1.48
truthvaltrue:condlow	1.49

Get predictions:

```
predict(glm1,  
        re.form = NA,  
        type = "resp")
```



## Data simulation for Study 2

```
n <- 200      # try different values


# Create data frame
dat <- data.frame(
  id          = factor(rep(1:n, each = 40)),
  item        = factor(paste(rep(1:5, each = 40), 1:40, sep = ":")),
  condition    = factor(rep(c("high", "low"), each = nitem)),
  truthvalue   = factor(rep(c("true", "false"), c(30, 10)))
)

# Do some checks
xtabs( ~ id + item, dat)
xtabs( ~ condition + truthvalue, dat)
xtabs( ~ condition + truthvalue + id, dat)
```

## Effect size of interest

How do I find suitable numbers for the model parameters?

Let us simulate some data with minimal error and plot the results.



	"true"
(Intercept)	0.00
truthvaltrue	0.05
condlow	-0.40
truthvaltrue:condlow	0.40

## Data simulation for Study 2

```
ran <- c("id.(Intercept)"    = 0.6,  
        "item.(Intercept)"  = 0.8)  
fix <- c("(Intercept)"      = 0,  
        "truthvaluetrue"    = 0.05,  
        "conditionlow"      = -0.4,  
        "truthvaluetrue:conditionlow" = 0.4)  
  
# Simulate data  
sim <- simulate( ~ truthvalue * condition + (1|id) + (1|item),  
                newdata      = dat,  
                newparams    = list(beta = fix, theta = ran),  
                family       = binomial,  
                nsim         = 20) # should be at least 400!
```

## Power simulation

```
pval <- numeric(ncol(sim))


for (i in seq_len(ncol(sim))) {
  m1 <- glmer(sim[, i] ~ truthvalue * condition +
              (1|id) + (1|item), dat, family = binomial)
  pval[i] <- summary(m1)$coef["truthvaluetrue:conditionlow", "Pr(>|z|)"]
}

# Power
mean(pval < 0.05)
hist(pval)
```

## Parameter recovery

```
par_rev <- replicate(400, {  
  sim <- simulate( ~ truthvalue*condition + (1|id) + (1|item),  
    newdata = dat,  
    newparams = list(beta = c(0, 0.01, -0.2, 0.2),  
      theta = c(0.5, 0.5)),  
    family = binomial)[,1]  
  glmer(sim ~ truthvalue * condition + (1|id) + (1|item),  
    data = dat, family = binomial)  
})  
mean(sapply(par_rev, fixef)[1,]); mean(sapply(par_rev, fixef)[2,])  
mean(sapply(par_rev, fixef)[3,]); mean(sapply(par_rev, fixef)[4,])  
  
mean(sqrt(sapply(par_rev, function(x) unlist(VarCorr(x)))[1,]))  
mean(sqrt(sapply(par_rev, function(x) unlist(VarCorr(x)))[2,]))
```

## Estimation of inflated effect





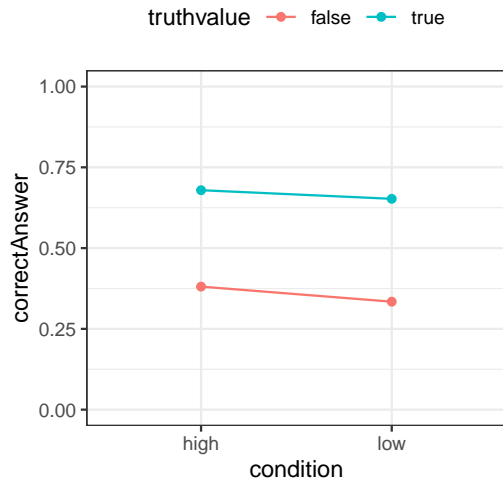
## Estimation of inflated effect

```
# Parameter estimates and confidence intervals
int <- sapply(par_rev,
  function(x) fixef(x)["truthvalue:conditionlow"])
ci <- sapply(par_rev,
  function(x) confint(x, method = "Wald")["truthvalue:conditionlow",])
# p values
p <- sapply(par_rev,
  function(x) summary(x)$coef["truthvalue:conditionlow", "Pr(>|z|)"])

# Power
mean(p < 0.05)
hist(p)
# Inflation of effect
summary(int[p < 0.05])
```

## Results of Study 2

	est	se	z	p
(Intercept)	-0.53	0.08	-6.66	0.00
truthvaltrue	1.34	0.07	18.72	0.00
condlow	-0.21	0.11	-1.90	0.06
truthvaltrue:condlow	0.08	0.10	0.74	0.46



## References

Wickelmaier, F. (2022). Simulating the power of statistical tests: A collection of R examples. *ArXiv*. <https://arxiv.org/abs/2110.09836>